

PartIR

Declarative abstractions for tensor program partitioning

Dimitrios Vytiniotis
dvytin@google.com

Dominik Grewe, Michael Schaarschmidt, James Molloy, Dan Belov

DeepMind

Adam Paszke, Dougal Maclaurin, Nicolas Vasilache

Google Brain

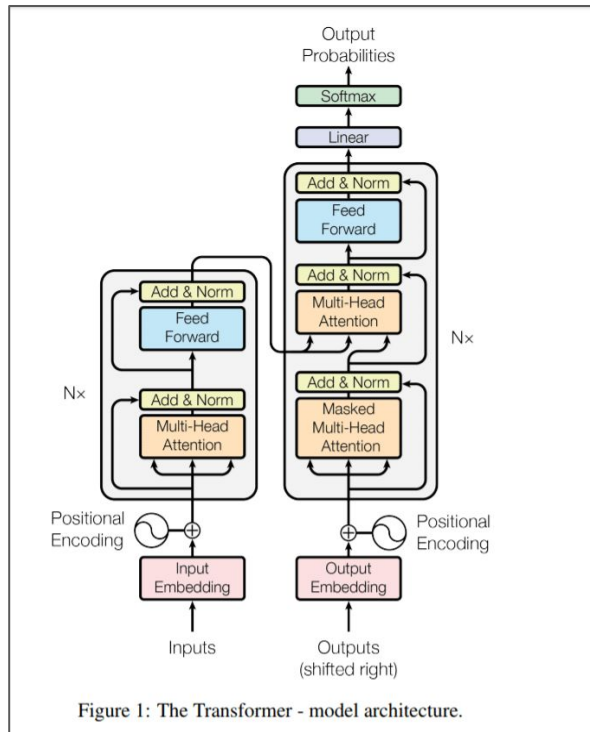
PPDP'20, September 2020



An invited talk about work-in-progress?!?

What is a machine learning model?

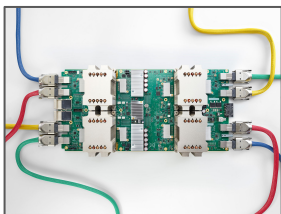
- Typically “beefy” multi-dimensional array programs: *tensor programs*
- Accept inputs and parameters:
 - Wild variation on input batch size, input spatial dimensions, and parameter size
- Parameters + compute grow:
 - GPT-3 is a **175B param** transformer model
- Ever-growing body of work to reduce parameters and compute:
 - Exploit sparsity [not the topic of this talk]
 - Novel ways to *partition these tensor programs* to exploit parallelism and make them fit memory constraints



Machine learning hardware: from tensor cores to systems

Cannot run a model with 175B parameters on a single machine!

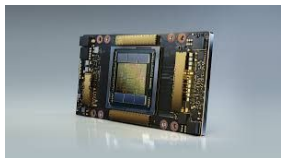
Accelerator trend for scaling: **chips** => **systems of multiple chips** => **data-centers**



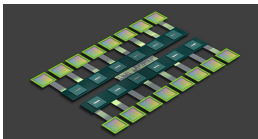
TPUs: few tensor
cores, ~10-100GBs of
accelerator memory



TPU pods: thousands of TPU
“cores” in custom interconnects



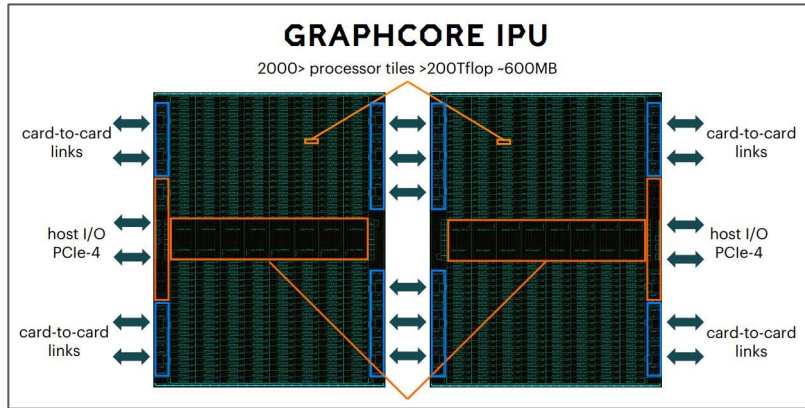
NVIDIA A100 GPU



nvswitch + nvlink for
data-center scale

Accelerator scaling trend: lots of wimpy cores

New and cheaper ways to scale compute and memory than monolithic MXU cores



The world's largest chip

46,225 mm² chip
56x larger than the biggest GPU ever made

400,000 cores
78x more cores

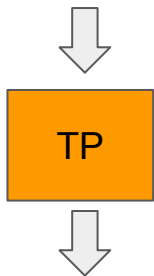
18 GB on-chip SRAM
3000x more on-chip memory

100 Pb/s interconnect
33,000x more bandwidth

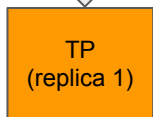
cerebras wse

Multiple ways to partition a tensor program

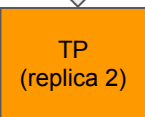
tensor<batch x size>



tensor<batch/2 x size>

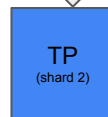
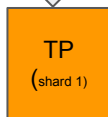


tensor<batch/2 x size>



Data Parallelism

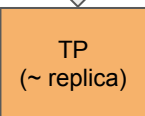
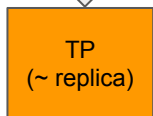
tensor<batch x size> tensor<batch x size>



Param sharding (v1)

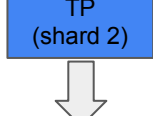
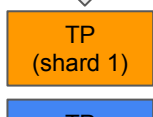
- Partitioning a model is a *program transformation*
- Often need to additionally mix-and match different forms in the same model
- **Beyond Data and Model Parallelism for Deep Neural Networks**, SysML'19

tensor<batch x size/2> tensor<batch x size/2>

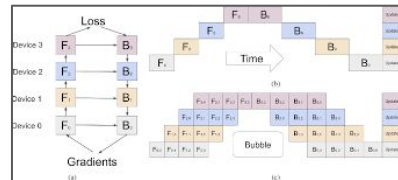


Input (feature) parallelism

tensor<batch x size>



Param sharding (v2)



Pipelining variant (gPipe)

Sounds simple? Still SW in a **BAD** place for partitioning

- We use data parallelism to reduce input (activation) sizes and scale the available compute flops, *assuming model parameters fit on available cores*
 - OK for TPUs and small models (~ 8-32GB/core), but new hardware may come with few GBs per device (e.g. Graphcore IPU). How does one map a large transformer model?
- We bake in system assumptions in library code
 - Typically batch parallelism, no model parallelism
- Engineering **years** on specialized partitions (papers: [Megatron](#), [gPipe](#)).
 - **New partitioning strategy ~> new research paper!**
- System design: manual back-of-the-envelope-calculations for mapping models on new systems. Hard to support a mixture of partitioning strategies

THIS IS MADNESS!!!!

PartIR aspiration: generic partitioning compiler infra

PartIR offers:

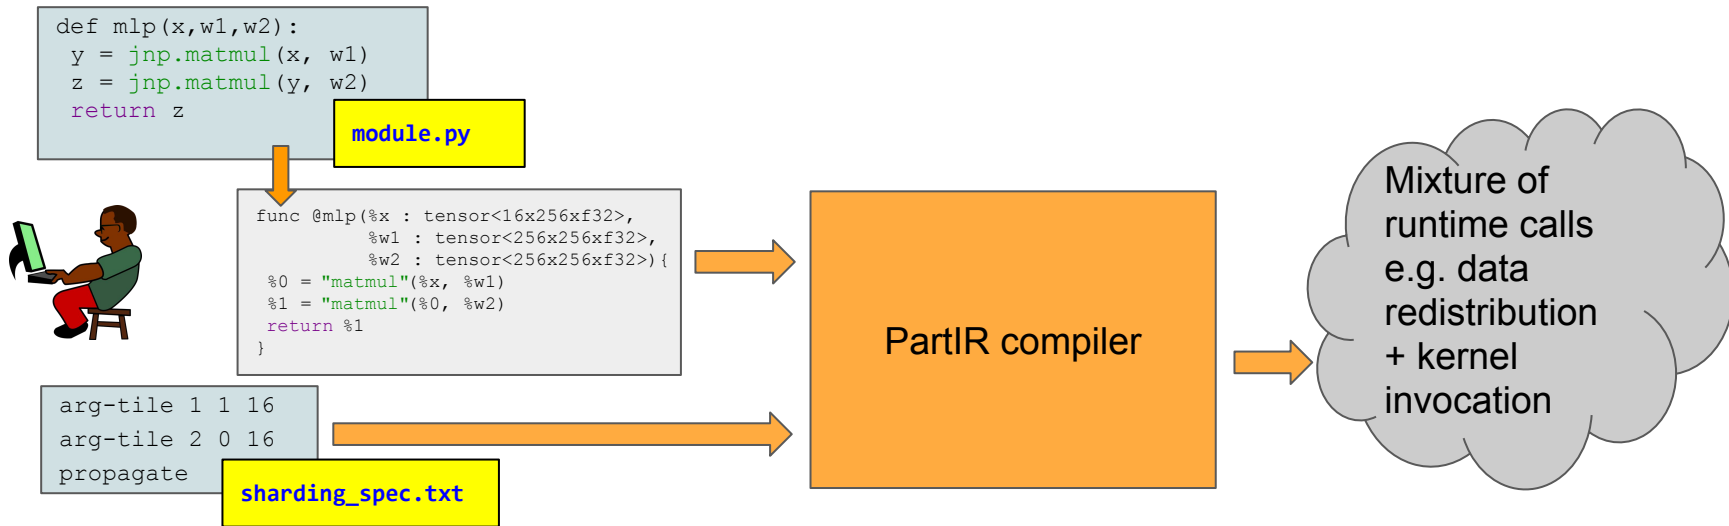
- Declarative **abstractions** to express partitioned models on accelerator systems
- Declarative **transformation rules for partitioning** actions and fusion
- Vision: driven by **sharding annotations**, or **interactive tactics**, or **search + RL** (non-exclusive)

Other features:

- POC lowering pass to dataflow graphs of SPMD ops (beyond a single SPMD op)
- Solves the “tiling-on-tensors” problem i.e. a functional specification of tiling
- Useful to explain semantics of related efforts (e.g. XLA sharding propagation, Mesh TF); eventually may increase expressivity (e.g. dataflow and control flow of SPMD ops, pipeline parallelism)
- Close relatives: [Dex](#), [Linalg](#), [F-smooth](#) [ICFP’19]
- Relies on some existing tensor op backend compiler and runtime

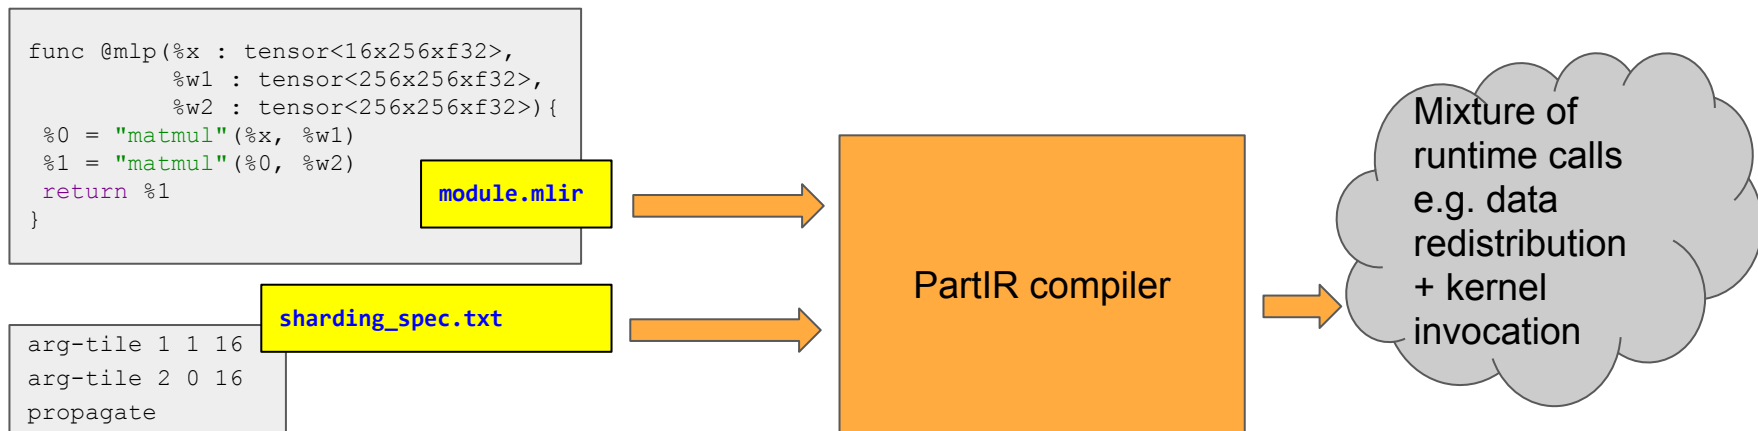
PartIR value; for ML researchers

Reusable math-focused (vs. *systems-focused*) ML libraries



PartIR value; for compiler infrastructure teams

Reusable partitioning compiler infra; across tensor dialects, different systems of accelerators

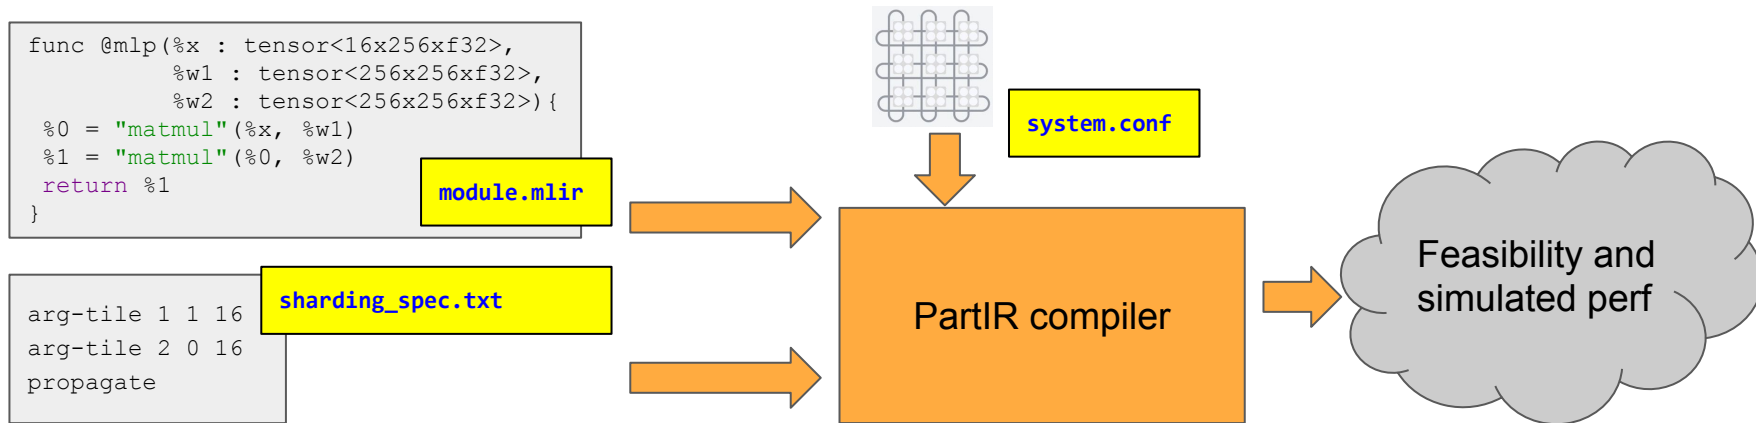


Do not invent a new IR from scratch: we have mature tensor IRs for different needs (e.g. XLA, LinAlg), need common infra to partition across accelerator systems. PartIR is just a generic partitioning framework on top: *it relies on a backend compiler and a runtime, for a given target (e.g. XLA for TPUs)*

PartIR value; to speed up systems design

Systems Design questions: what if 2x memory/core? 2x cores with $\frac{1}{2}$ mem? 2x faster mesh? 2x FLOPS/sec/core?

How can we quickly experiment with partitioning and mapping important models in different ways? At the moment it's months of engineering time but a ton of boilerplate.



A technical tour of PartIR

PartIR sits **on top of some other tensor dialect**, e.g. XLA:

- Using the MLIR (multi-level IR, <https://mlir.llvm.org/>) framework this is technically easy to express
- MLIR is infrastructure for representing multiple IRs (*dialects*). Allows programs to span operators from different dialects and features generic binding trees facilities, generic representations of modules, functions, SSA blocks, rewriting APIs, an extensible type system and more
- XLA is a dialect in MLIR and PartIR simply adds a few constructs as a separate dialect

PartIR constructs:

- A **range type** (partir.range) whose values represent contiguous non-overlapping intervals
- 3 operators specific to partitioning:
 - An op to **slice a dimension** (partir.slice)
 - An op to **tile a dimension** (partir.tile)
 - An op to do a **tiled reduction** (partir.sum)

Not discussed today: a handful of extra operators that allow tensor *building* (as opposed to *partitioning*): out of scope for today's talk, but very useful for e.g. Dex. Such operators were first introduced in the F-smooth paper.

PartIR range type and its values: $\text{range}\langle n, w \rangle$

A range type $\text{range}\langle n, w \rangle$ denotes the set of contiguous non-overlapping intervals of a range n , each interval being of width w .

- Assume $(x:\text{range}\langle 16, 8 \rangle)$. Then x is one of: $[0..8)$, $[8..16)$
- Non-overlapped and **perfectly divisible**, for now. ($n \text{ `mod` } w == 0$)
- We can nest ranges in a monoid fashion (will skip for this talk)

PartIR constructs: slicing a dimension with a range value

Key idea: use range values directly in slicing

slice d x[r]

A dimension
index (attribute)

A tensor-typed
value

A range-typed
value

```
x : tensor<64x32x64xf32>
```

```
r : range<64,4>
```

```
slice 0 x[r] : tensor<4x32x64xf32>
```

```
slice 1 x[r] : TYPE-ERROR
```

```
slice 2 x[r] : tensor<64x32x4xf32>
```

PartIR constructs: tiling a dimension with a range

PartIR introduces a higher-order loop-like expression for this:

A constant integer, which
dimension to slice, here d=1

```
y = tile d (\(r : range<n,w>) -> expr)
```

y : tensor<32x~~n~~x16xf32>

expr : tensor<32x~~w~~x16xf32>

Semantics: a generator expression for the slices of a bigger array. It can be given either parallel or sequential semantics (depends on lowering)

PartIR constructs: reductions

A similar higher-order operator:

$y = \text{sum } (\backslash(r : \text{range}\langle n, w \rangle) \rightarrow \text{expr})$

$y : \text{tensor}\langle 32 \times 16 \times f32 \rangle$

$x : \text{tensor}\langle 32 \times 16 \times f32 \rangle$

Semantics: sum together all the $\langle 32 \times 16 \times f32 \rangle$ chunks to a single tensor of the same shape. Can be implemented with all-reduce in a distributed setting (see later)

Partitioning = progressive application of rewrite rules

```
x : tensor<nxmxf32>, y : tensor<mxoxf32>
matmul-tile-0:
  matmul(x, y) <~~> tile 0 (\r:range(n,nw) -> matmul(slice 0 x[r], y))
matmul-tile-1:
  matmul(x, y) <~~> tile 1 (\r:range(o,ow) -> matmul(x, slice 1 y[r]))
matmul-sum:
  matmul(x, y) <~~>
    sum (\r:range(m,mw) -> matmul(slice 1 x[r], slice 0 y[r]))
```

- Rewrite rules preserve types and semantics
- Each tensor operator from our base dialect is equipped with annotations that inform how each dimension could be partitioned, and how this propagates to argument slicing.
- **Opportunity**: given the mathematical definition of an op or even a sub-program search for valid rewrites + validate through a theorem prover. Reminiscent of TASO [SOSP'19])

Producer-consumer fusion = when slice met tile

```
s : range<64,32>
```

```
fuse:
```

```
  slice 0 (tile 0 (\r:range<64,32> -> expr)[s] ~~> expr{s/r})
```

Analogous to the build-slice fusion rule from F-smooth paper [ICFP'19]

Slicing function arguments or values (aka “dumb tiling”)

```
x : tensor<64x32xf32>
```

```
x <~~> tile 0 (\r:range(64,4) -> slice 0 x[r])
```

```
x <~~> tile 1 (\r:range(32,4) -> slice 1 x[r])
```

Propagation tactics - push slicing in

```
slice 0 (matmul(x, y))[r] ~~> matmul(slice 0 x[r], y)
```

```
slice 1 (matmul(x, y))[r] ~~> matmul(x, slice 1 y[r])
```

Intuition: do not compute a big matmul only to take a slice out of it! Instead, directly compute the slice. Also note: these are derived rules from fusion + the specialized matmul ops, but useful on their own!

Propagation tactics - pulling tiling **out**

```
matmul(tile 0 (\r:range(n,nw) -> expr), y) ~~>  
      tile 0 (\r:range(n,nw) -> matmul(expr, y))  
  
matmul(x, tile 1 (\r:range(o,nw) -> expr)) ~~>  
      tile 1 (\r:range(o,ow) -> matmul(x, expr))
```

Intuition: try to expose looping constructs at the top-level

But needs to be done carefully: we are pulling “y” **into** the tiling construct, hence likely to increase the local memory requirements on every compute core!

A space of user-controlled actions + propagation

```
func @mlp(%x: tensor<16x256xf32>,  
          %w: tensor<256x256xf32>,  
          %u: tensor<256x256xf32>)  
-> tensor<16x256xf32> {  
  %0 = matmul(%x, %w)  
  %1 = matmul(%0, %u)  
  return %1  
}
```

arg-tile 0 0 8.



```
func @mlp(%x: tensor<16x256xf32>, %w: tensor<256x256xf32>,  
          %u: tensor<256x256xf32>) -> tensor<16x256xf32> {  
  %0 = partir.tile 0 (%r : !partir.range<16,8>) {  
    %3 = partir.slice 0 %x[%r]  
    partir.yield %3  
  }  
  %1 = matmul(%0, %w)  
  %2 = matmul(%1, %u)  
  return %2  
}
```

Data parallelism

propagate.



```
func @mlp(%x: tensor<16x256xf32>, %w: tensor<256x256xf32>, %u: tensor<256x256xf32>)  
-> tensor<16x256xf32> {  
  %0 = partir.tile 0 (%r : !partir.range<16,8>) {  
    %1 = partir.slice 0 %x[%r]  
    %2 = matmul(%1, %w)  
    %3 = matmul(%2, %u)  
    partir.yield %3  
  }  
  return %0  
}
```

**arg-tile 1 1 128;
arg-tile 2 0 128;
propagate.**



```
func @mlp(%x: tensor<>, %w: tensor<256x256xf32>,  
          %u: tensor<256x256xf32>) -> tensor<16x256xf32> {  
  %0 = partir.tile 0 (%r : !partir.range<16,8>) {  
    %1 = partir.slice 0 %x[%r]  
    %2 = partir.sum (%s : !partir.range<256,128>) {  
      %3 = partir.slice 1 %w[%s]  
      %4 = matmul(%1, %3)  
      %5 = partir.slice 0 %u[%s]  
      %6 = matmul(%4, %5)  
      partir.yield %6  
    }  
    partir.yield %2  
  }  
  return %0  
}
```

**Data parallelism
+
parameter sharding**

Lowering PartIR

Since many accelerators offer facilities for:

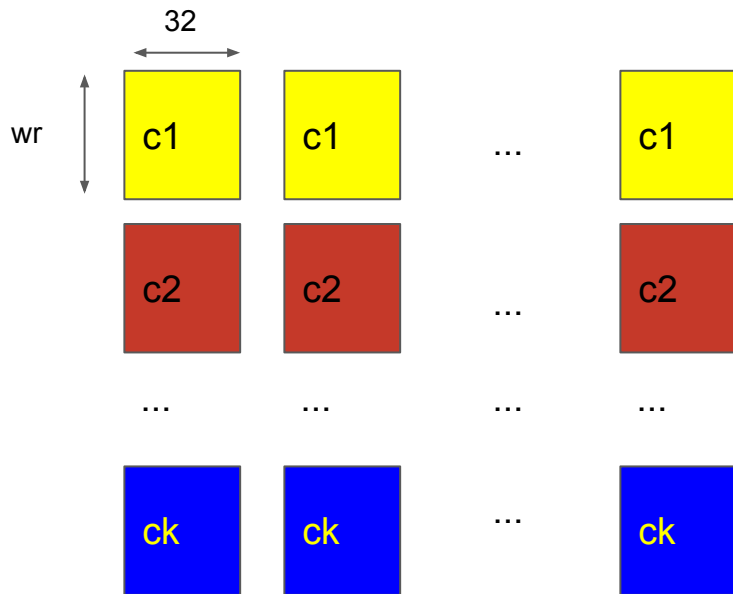
- data transfers between devices and re-distribution
- SPMD-style ops across a number of available cores

we sketch out a translation of PartIR programs to such an SPMD dialect, which we call PartIR:SPMD. This entails several steps:

- Introducing the type system of PartIR:SPMD
- Introducing the ops of PartIR:SPMD
- Introducing progressive lowering to PartIR:SPMD

PartIR:SPMD Distributed tensor types

`distributed_tensor<(r:range(nr,wr), s:range(ns,ws)) -> [r,32]>`



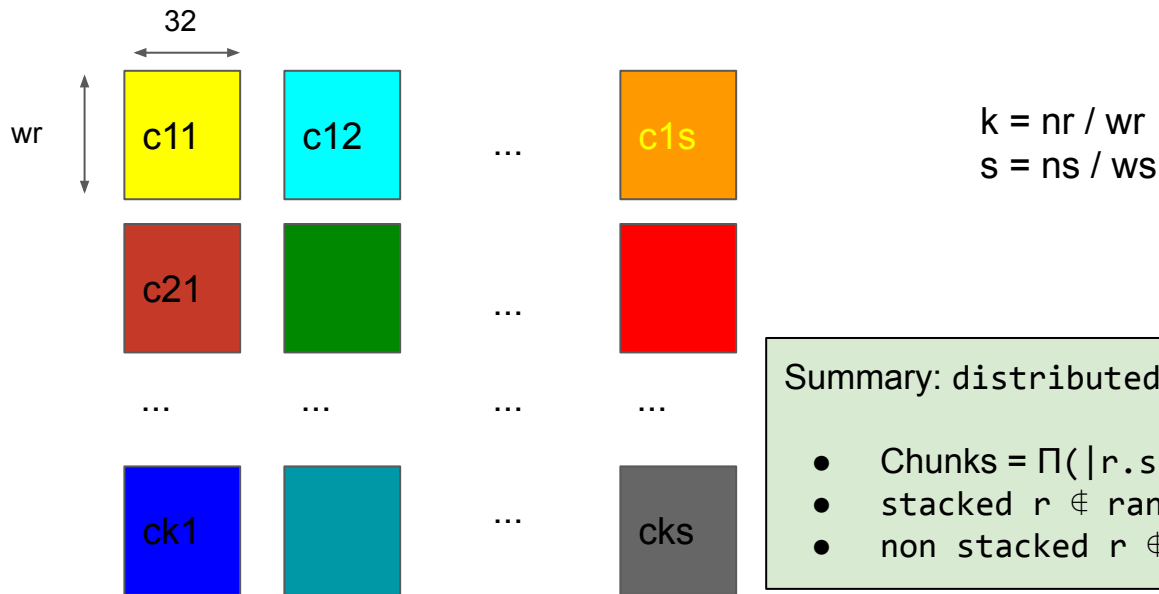
$$k = nr / wr$$

$$\text{horizontal replication} = ns/ws$$

binding
occurrences

PartIR:SPMD Stacked tensor types

`distributed_tensor<(r:range(nr,wr), @stacked s:range(ns,ws)) -> [r,32]>`



Summary: `distributed_tensor<(rs) -> shape>`

- $\text{Chunks} = \prod (|r.\text{size}/r.\text{tilesize}| \text{ for } r \text{ in } rs)$
- stacked $r \notin \text{rangeVars}(\text{shape}) \Rightarrow \text{stacking}$
- non stacked $r \notin \text{rangeVars}(\text{shape}) \Rightarrow \text{replication}$

Lowering step 1: introduce SPMD op + lift free variables

```
%x : tensor<32x8xf32>  
%y : tensor<8x16xf32>  
%0 = tile 0 (%r:range(32,16) -> matmul(slice 0 %x[%r], %y)
```

Lift free variables %x and %y via replication, naively



```
%y1 = distribute(%y) : distributed_tensor<(%r:range(32,16)) -> [8,16]>  
%x1 = distribute(%x) : distributed_tensor<(%r:range(32,16)) -> [32,8]>  
%1 = spmd(%x1, %y1) (%r:range(32,16),  
                    %x_arg : tensor<32x8xf32>,  
                    %y_arg : tensor<8x16xf32>) {  
    %2 = matmul(slice 0 %x_arg[%r], %y_arg);  
    yield %2  
} : distributed_tensor<(@stacked %r:range(32,16)) -> [16,16]>  
%2 = tile_stacked_tensor 0 (%1) : distributed_tensor<(r:range(32,16)) -> [r, 16]>  
%3 = undistribute(%2) : tensor<32x16xf32>
```

Introduce a generic SPMD op on multiple arguments

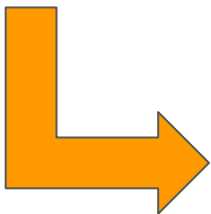
SPMD op returns a “stacked” tensor

Tile (i.e. rearrange stacked tensor into distributed tensor

Explicitly un-distribute to preserve original type

Lowering step 2: transform replication to distribution

```
%y1 = distribute(%y) : distributed_tensor<(%r:range(32,16)) -> [8,16]>
%x1 = distribute(%x) : distributed_tensor<(%r:range(32,16)) -> [32,8]>
%1 = spmd_op(%x1, %y1) (%r:range(32,16),
    %x_arg : tensor<32x8xf32>,
    %y_arg : tensor<8x16xf32>) {
    %2 = matmul(slice 0,%x_arg[%r],%y_arg);
    yield %2
} : distributed_tensor<(@stacked %r:range(32,16)) -> [16,16]>
%2 = tile_stacked_tensor 0 (%1) : distributed_tensor<(r:range(32,16)) -> [r, 16]>
%3 = undistribute(%2) : tensor<32x16xf32>
```



```
%y1 = distribute(%y) : distributed_tensor<(%r:range(32,16)) -> [8,16]>
%x1 = distribute(%x) : distributed_tensor<(%r:range(32,16)) -> [32,8]>
%x1 = distribute(%x) : distributed_tensor<(%r0:range(32,16)) -> [%r0,8]>
%1 = spmd_op(%x1, %y1) (%r:range(32,16),
    %x_arg : tensor<16x8xf32>,
    %y_arg : tensor<8x16xf32>) {
    %2 = matmul(%x_arg, %y_arg);
    yield %2
} : distributed_tensor<(@stacked %r:range(32,16)) -> [16,16]>
%2 = tile_stacked_tensor 0 (%1) : distributed_tensor<(r:range(32,16)) -> [r, 16]>
%3 = undistribute(%2) : tensor<32x16xf32>
```

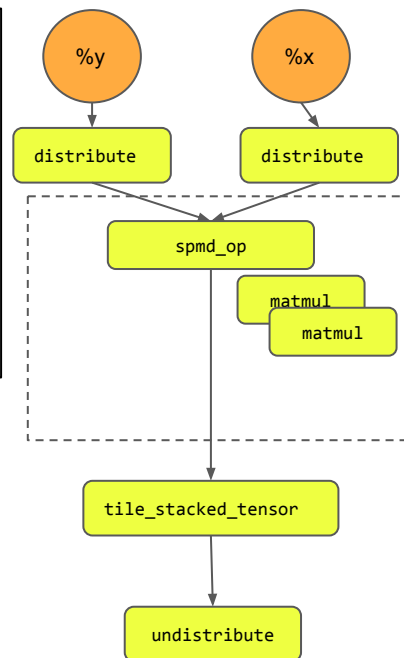
More lowering details (known and unknown)

- For translating `partir.sum` we introduce `spmv + partir.all_reduce` op
- Nested `partir.tile` and `partir.sum`
- Non-perfectly nested code inside `partir.tile/partir.sum`
 - Need to choose between creating distributed versions of intermediates, or inlining in inner loops and replicating computation. Expose option to programmers
- Fusion of distribution operators:
 - `undistribute(distribute[τ](%x) ~> %x`
 - `distribute[τ](undistribute(%x) ~> %x`
 `when type(%x) == τ`
 - `distribute[τ](undistribute(%x) ~> redistribute %x`
 `when globalType(type(%x) == globalType(τ)`

Execution

```
%y1 = distribute(%y) : distributed_tensor<(%r:range(32,16)) -> [8,16]>
%x1 = distribute(%x) : distributed_tensor<(%r0:range(32,16)) -> [%r0,8]>
%1 = spmd_op(%x1, %y1) (%r:range(32,16),
    %x_arg : tensor<16x8xf32>,
    %y_arg : tensor<8x16xf32>) {
    %2 = matmul(%x_arg, %y_arg);
    yield %2
} : distributed_tensor<(@stacked %r:range(32,16)) -> [16,16]>
%2 = tile_stacked_tensor 0 (%1) : distributed_tensor<(r:range(32,16)) -> [r, 16]>
%3 = undistribute(%2) : tensor<32x16xf32>
```

- Compile all bodies of spmd_op using the backend compiler
- Literally execute the data-flow graph using the backend runtime (must support the redistribution operators and SPMD execution)



Range widths that do not divide the dimension?

Proposal: use dependent types

- Semantics: $\text{Range}(16, 5) = [0..4), [5..10), [10..15), [15, 16)$
- Introduce dependent types:
 - `slice d %x[%r] : <n1, n2, ..., %r, ..., n_k>`
- Type `partir.tile` using dependent types:

I.e. let the last interval have a smaller width

BEFORE

`y : tensor<32xnx16xf32>`

`x : tensor<32xwx16xf32>`

```
%y = tile d (%r : range<n,w>) { ... stuff ... ; yield %x }
```

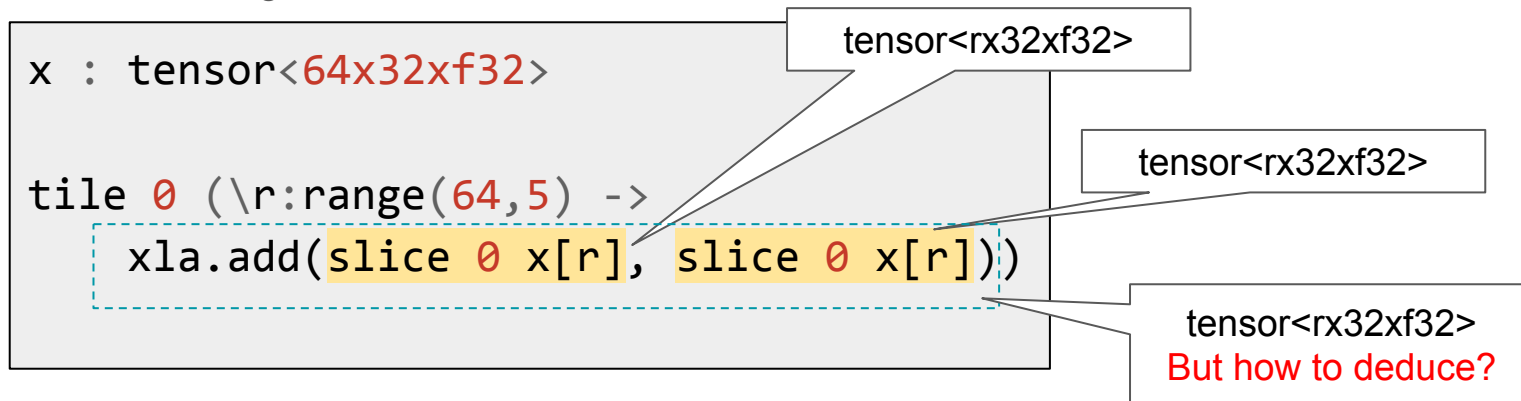
`y : tensor<32xnx16xf32>`

`x : tensor<32xrx16xf32>`

NOW

More expressive (dependent) types annoyance

- XLA dialect only has static shape: **not dependent types on range values!**
- How to type check a PartIR + XLA program with dependent types?
- Solution: specialization
 - `range<32,5>` has values with in-effect 2 widths: 5 and 2
 - type-check existing XLA ops for the two different assignments of widths
 - can lead to exponential blow-up for deeply nested tiles/sum, but we don't really expect a nesting level of > 2-3



An active research & engineering effort in DeepMind

- Testing strengths and weaknesses of PartIR on a set of mission-important models (XLA generated from Tensorflow or Jax high-level programs)
- Introduced “tiling specifications” for a substantial subset of XLA ops
- Working through lowering and execution for TPU/GPU environments
- Started work on exposing the state of a partir program and an action space as a reinforcement-learning environment to tap onto the big pool of ML researchers in DeepMind. Example actions:
 - Application of a rewrite rule
 - Navigation
 - Effective inlining of a tile op inside another
 - Pushing tiling out
- Work towards approximate cost models for evaluating solutions (e.g. minimize data redistribution, subject to memory fit) and learnt cost models from data (Ithema1, [ICML'19])
- Identifying architectural modifications to MLIR pass and rewriting APIs needed for RL

What we do not know: how much declarative is enough?

- How to best express even higher-level abstractions, most notably: pipelining
- Input (“feature”) partitioning when we can’t exactly partition in the boundaries:
 - halo exchanges and convolutions require extra data to be communicated
- Is a tile-based IR on top of a tensor IR already too low level? Do we need higher-level combinators -- for instance “map over data” or “map over parameters” or “pipeline”?
- How to enhance numpy-style tensor programs with PartIR concepts:
 - Expose PartIR constructs directly, let people program with those (not portable code)
 - Numpy + interactive environment to partition a program (nice but also niche!)
 - Numpy + sharding annotations on variables (a bit better but limits library usability)
 - Numpy + sharding annotations on inputs and outputs only + ML/search for intermediates?

Thank you!

- We desperately need abstractions for partitioning in our compiler stacks
 - Types, Operators, Transformations
- PartIR offers a principled approach to partitioning via semantics preserving sequences of really simple transforms, rooted in deforestation and fusion ideas from declarative programming
- Great value for understanding the partitioning problem and having an executable semantics of partitioned programs
- But many things we do not know and need to learn; e.g. what has the HPC/MPI community been doing in this space?

The work is really just starting, keen to engage and collaborate!

Extra material

Notation convention: SSA vs. expressions

We use expression-based notation instead of SSA cause it's shorter.

```
tile d (\r:range<n,w> -> expr)
```

expr here can be thought of as the body of a region that yields a value.

```
tile 0 (\r:range(64,4) -> matmul(slice 0 x[r],y))
```

NOTATION

```
tile 0 (%r:range(64,4)) {  
    %0 = slice 0 %x[%r];  
    %1 = matmul(%0, %y);  
    yield %1 }
```

ACTUAL MLIR